

Growth Marketing Assignment

Rippling

This exercise will give you an idea of the role, and help us understand your skills and abilities.

We generally look to receive the assignment within 72 hours, but please don't hesitate to reach out if you need additional time.

If you have any questions, please direct them to Will (whunter@rippling.com). Please email Will and Megan (megan.ballard@rippling.com) once completed!

About Rippling

Rippling gives businesses one place to run HR, IT, and Finance. We support thousands of sellers across 30+ products, serve companies of all sizes, and compete across multiple categories (payroll, HRIS, spend management, IT, workforce management, and more). You can learn more at rippling.com feel free to leverage any publicly available material on our website to build your understanding of who we are and how we position.

The Challenge

Build a conversational AI agent that takes a competitor's name or domain as input and produces a structured analysis of a competitor's public marketing strategy and positioning. The agent should work for *any* competitor in our space; it should not be hardcoded to a single company.

This agent analyzes publicly available marketing materials only — ad libraries, public websites, social media posts, press releases. It is not meant to uncover proprietary information, product roadmaps, or anything behind a login.

**The goal is to understand how competitors market and position themselves using only publicly available information — not to reverse-engineer their product roadmap or internal strategy.*

What the agent should do:

1. **Research across multiple sources:** The agent should gather intelligence from a combination of:
 - Paid ad creative and messaging (Meta Ad Library, Google Ads Transparency Center)
 - Website positioning, product pages, and pricing
 - Social media presence and recent content themes
 - Recent product launches, press, or announcements
 - Any other sources you think are valuable
2. **Make decisions:** The agent should have a real agentic loop. Based on what it finds (or doesn't find), it should decide what's worth researching deeper, what to skip, and when to move on. Not every source will have useful data for every competitor.
3. **Be conversational:** The agent should be interactive. If it has clarifying questions ("Should I focus on their enterprise or SMB positioning?" or "I can't find active Meta ads want me to try LinkedIn instead?"), it should surface them to the user. It should handle follow-up requests ("Now dig deeper on their pricing" or "Run this again for Gusto").
5. **Produce structured output:** The final deliverable from the agent should include:
 - A **markdown brief** covering:
 - What messaging angles and themes they're running
 - How they position their product(s)
 - What's changed recently (new campaigns, new ICPs being targeted, messaging pivots)
 - **What positioning gaps or opportunities does this surface for Rippling's marketing?** What would you exploit if you were writing campaigns against them?
 - A **JSON file** with the structured data the agent collected (sources, extracted claims, confidence levels, timestamps)

You may use any tools, APIs, models, or frameworks you want. All research should use publicly available information only (ad libraries, public web pages, social posts, press). Do not scrape behind logins or access any non-public data.

Deliverables

1. **Code:** GitHub repo or zip file. We should be able to run the agent locally.
2. **Agent output:** The actual competitive brief the agent produced for at least one competitor of your choice.
3. **Loom video (5-10 min):** Walk us through
 - Your architecture decisions (why you structured the agent this way)
 - Your tool and model choices (what you used and why)
 - How the eval system works
 - A live demo of the agent running

What we're evaluating

- **Does the agent work?:** Can we point it at a competitor and get a useful, accurate picture of their marketing strategy?
- **Architecture quality:** Is it correctly structured as an agent (not a linear script)? Does it make good decisions about what to research and when?
- **Tool usage:** Does it use the right tools for the job? Does it handle tool failures gracefully?
- **Output quality:** Is the brief actually useful? Is the "relevance to Rippling" section insightful?
- **Decisions and trade-offs:** We care more about *why* you made choices than whether the code is perfect. The Loom is your chance to show your thinking.

We are not evaluating

- UI/frontend: (CLI is fine; a UI is bonus but not expected)
- Code style perfection: (readable and functional > beautifully formatted)
- Whether you use a specific framework or language